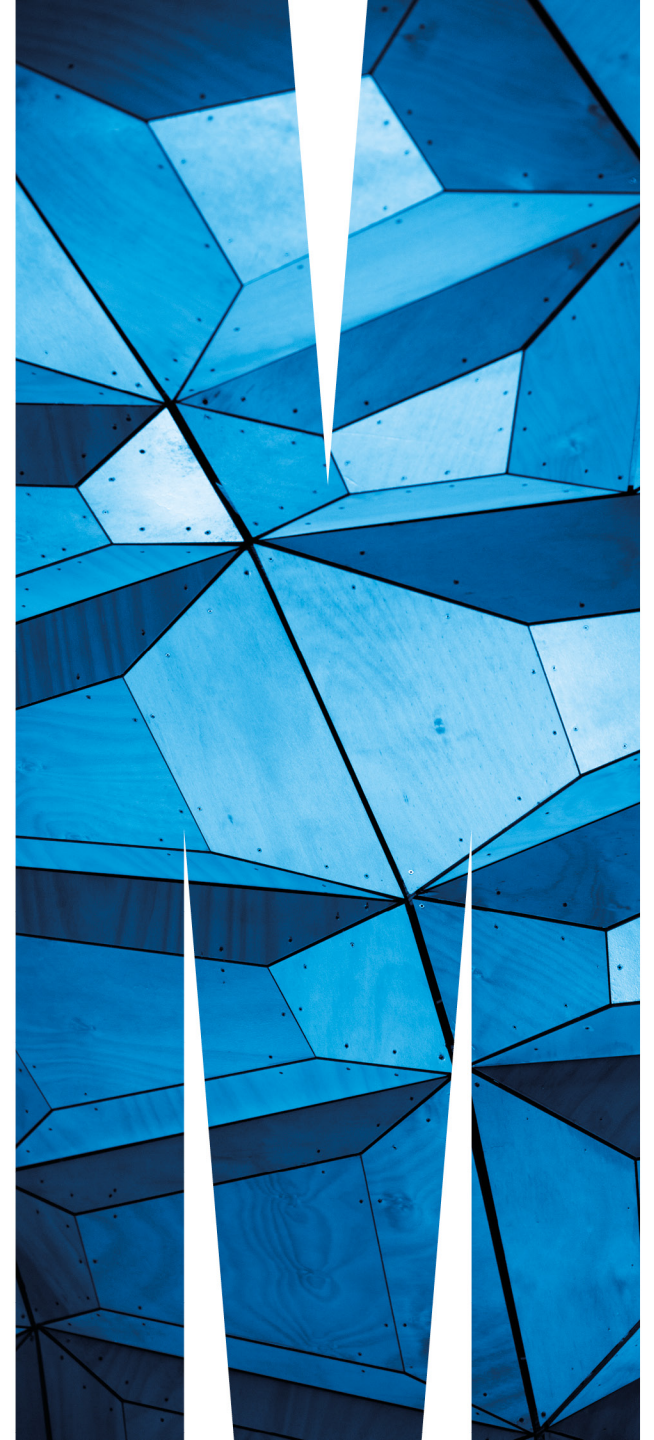


FIT2100 Semester 2 2017

Lecture 2: Operating System Overview (Reading: Stallings, Chapter 2)

Jojo Wong



Lecture 2: Learning Outcomes

- Upon the completion of this lecture, you should be able to:
 - Understand the objectives and functions of an operating system (OS)
 - Discuss the evolution of operating systems from simple batch systems to modern complex systems
 - Discuss the key design areas of modern operating systems
 - Understand the OS architecture of Unix/Linux, Android, and Windows



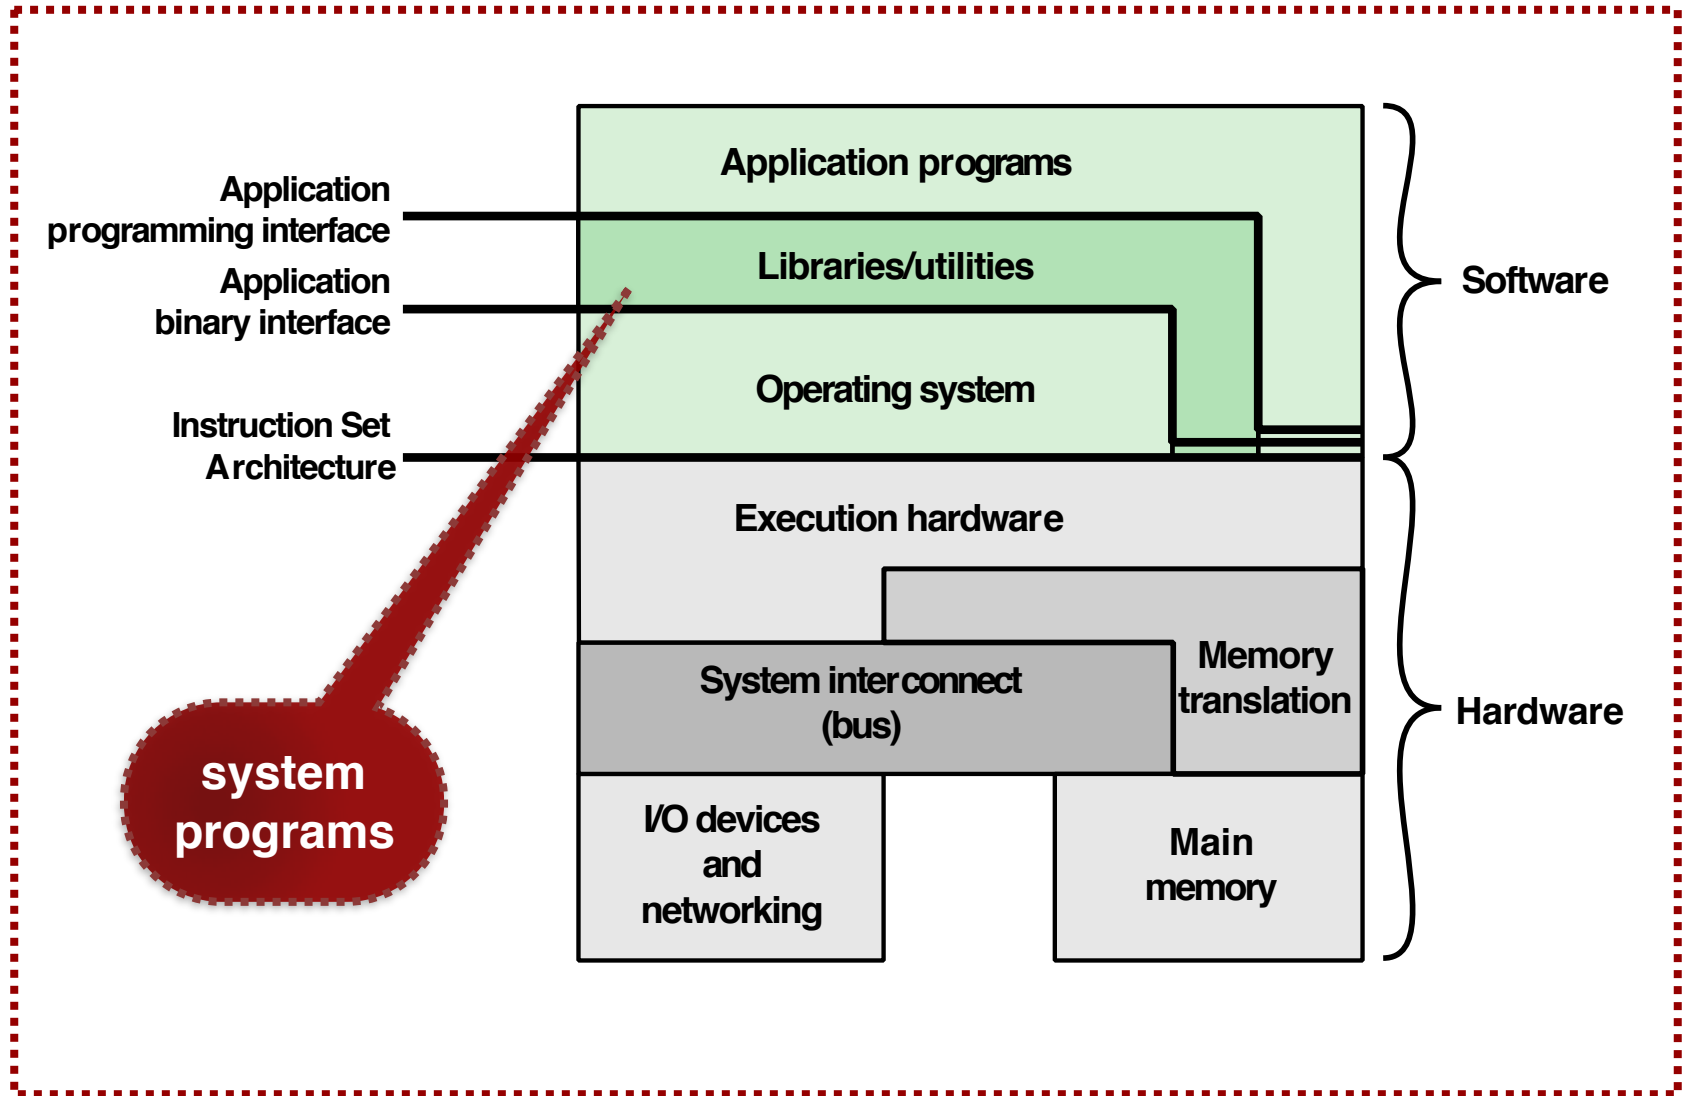
MONASH
University

What is an operating system?

Operating System

- A **program** that controls the execution of application programs
- An **interface** between applications and hardware
- Main **objectives** of an operating system:
 - Convenience
 - Efficiency
 - Ability to evolve

Computer Hardware and Software Structure



Operating System Services

- Program development
- Program execution
- Access I/O devices
- Controlled access to files
- System access
- Error detection and response
- Accounting

What is the role of an operating system?

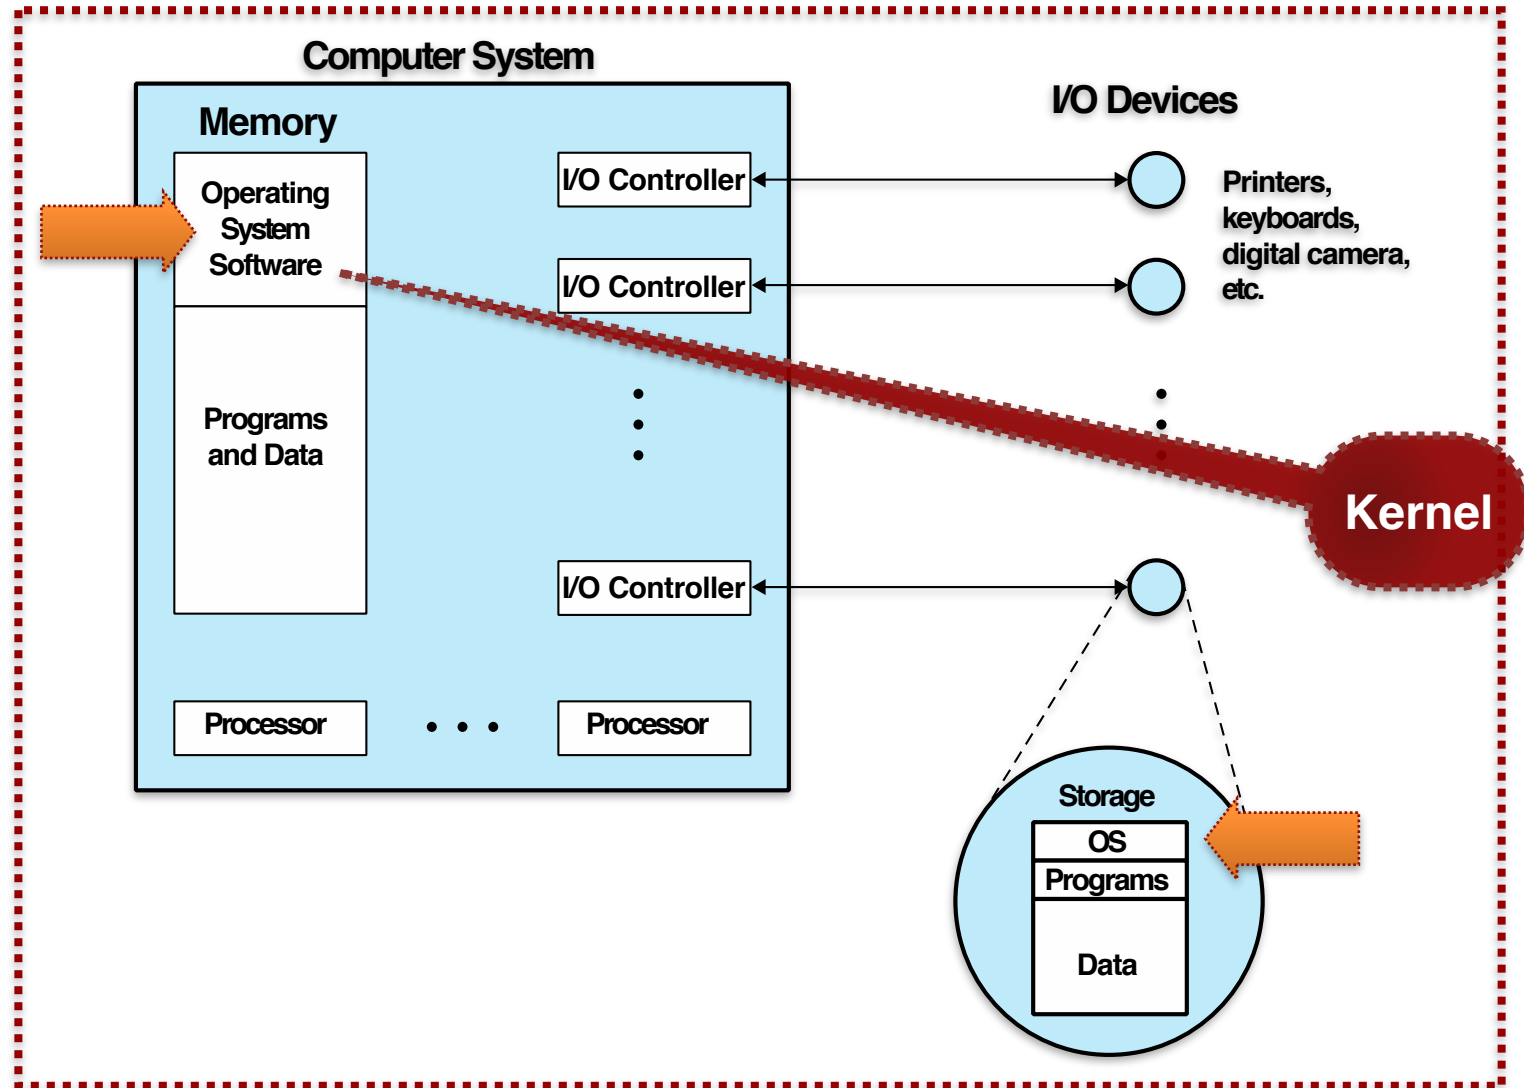
The Role of an OS

- A computer is a set of resources for the movement, storage, and processing of data
- The OS is responsible for managing these resources
- The OS is in control of the computer's basic functions

Operating System: As Software

- OS functions in the same way as ordinary computer software
- Program, or **suite of programs**, executed by the processor
- Frequently **relinquishes control** and must depend on the processor to allow it to regain control

Operating Systems: As Resource Manager



Operating System: Other Terms

- **Resource allocator:**
 - Manages and allocates resources to programs
- **Control program:**
 - Controls the execution of user programs and operation of I/O devices
- **Kernel:**
 - The one program running at all times (all else being application programs)



Unix/Linux
OS

The evolution of operating systems

The Evolution of Operating Systems

- A major OS will **evolve** over time for a number of reasons:

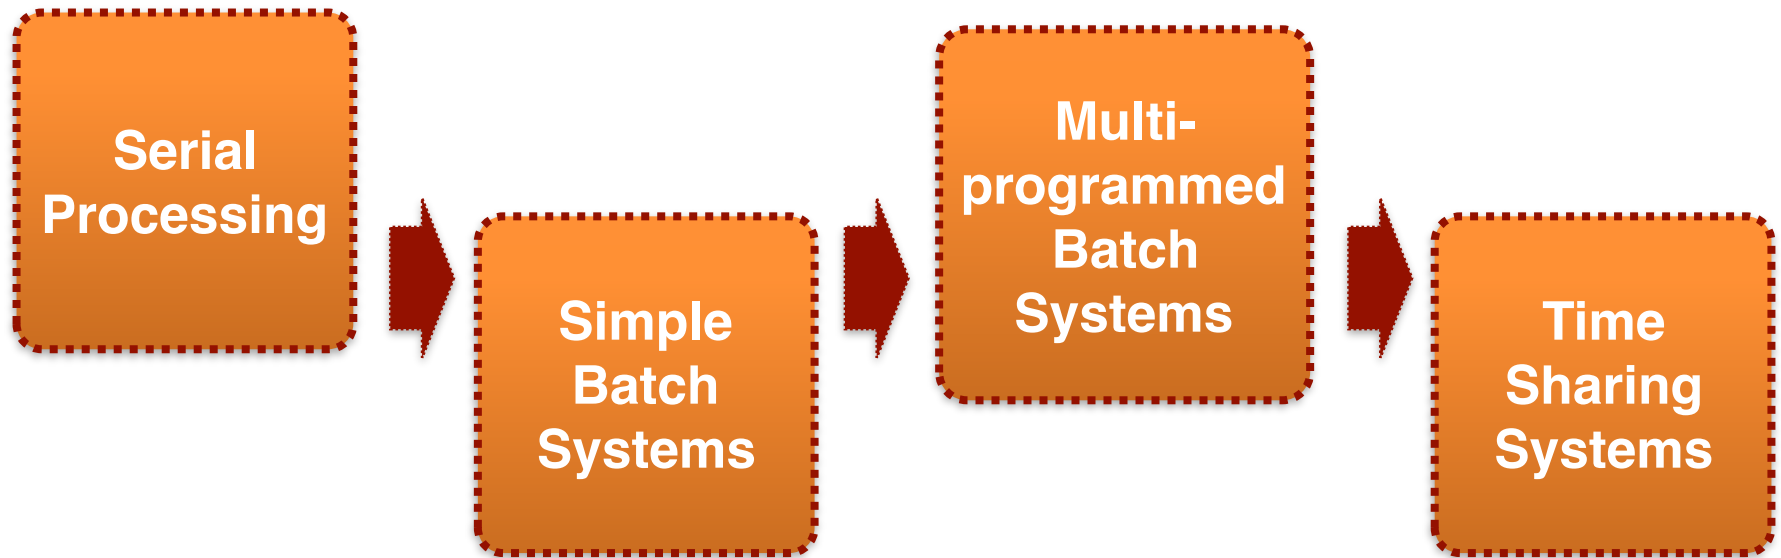
Hardware upgrades

New types of hardware

New services

Fixes

The Evolution of Operating Systems



Serial Processing

Earliest Computers

- No operating system
 - Programmers interacted directly with the computer hardware
- Computers ran from a console with display lights, toggle switches, some form of input device, and a printer
- Users have access to the computer in “series”

Problems


- Scheduling
 - Most installations used a hardcopy sign-up sheet to reserve computer time
 - Time allocations could run short or long, resulting in wasted computer time
- Setup time
 - A considerable amount of time was spent just on setting up the program (job) to run

Simple Batch Systems

- Early computers were very expensive
 - Important to **maximise processor utilisation**
- The concept of **monitor**:
 - Users no longer have direct access to the processor
 - Jobs are submitted to a computer operator who batches them together and places them on an input device
 - Program branches back to the monitor when finished

Simple Batch Systems: Job Scheduling

Special type of programming language (JCL) used to provide instructions to the monitor



What *compiler* to use?



What *data* to use?

Simple Batch Systems: Modes of Operation

User Mode

- User program executes in user mode
- Certain areas of memory are protected from user access
- Certain instructions may not be executed

Kernel Mode

- Monitor executes in kernel mode
- Privileged instructions may be executed
- Protected areas of memory may be accessed

Simple Batch Systems: Overhead

- Processor time alternates between execution of user programs and execution of the monitor
- **Sacrifices:**
 - Some main memory is given over to the monitor
 - Some processor time is consumed by the monitor
- Despite overhead, the simple batch system improves utilisation of the computer

Multiprogrammed Batch Systems

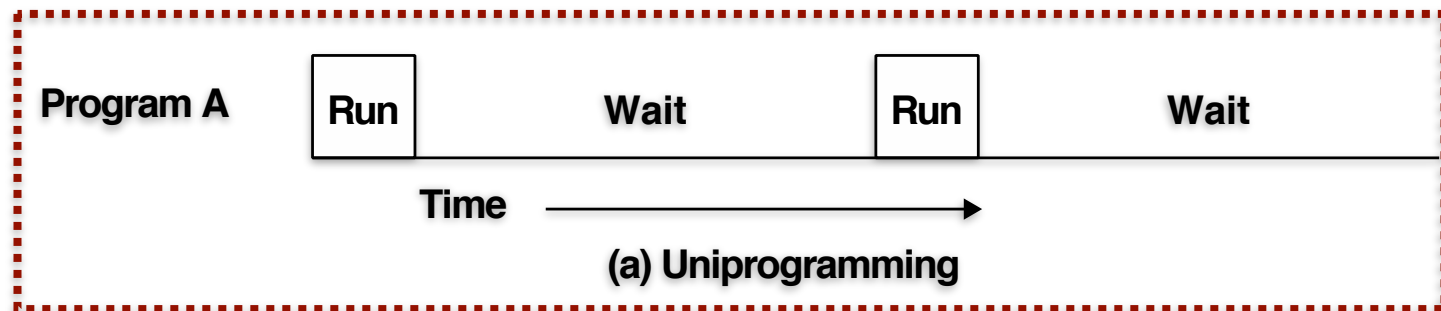
- Processor is often **idle**:
 - Even with automatic job sequencing
 - I/O devices are slow as compared to processor
- Example: system utilisation

Read one record from file	15 μ s
Execute 100 instructions	1 μ s
Write one record to file	<u>15 μs</u>
TOTAL	31 μ s

$$\text{Percent CPU Utilization} = \frac{1}{31} = 0.032 = 3.2\%$$

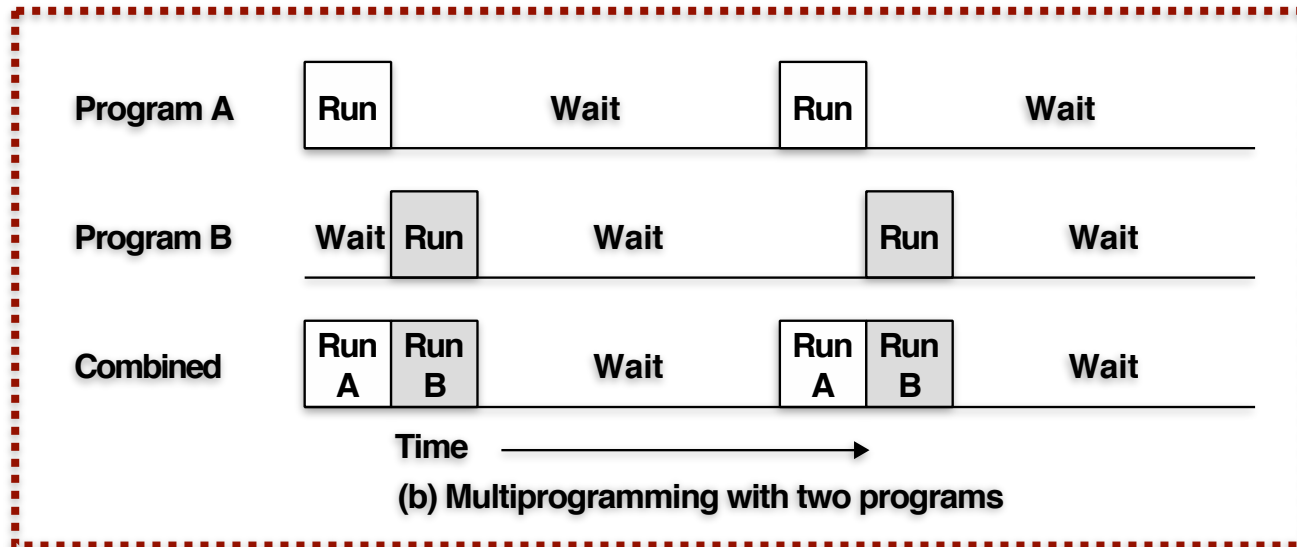
Uniprogramming

- The processor spends a certain amount of time executing, until it reaches an I/O instruction
- It must then wait until that I/O instruction concludes before proceeding



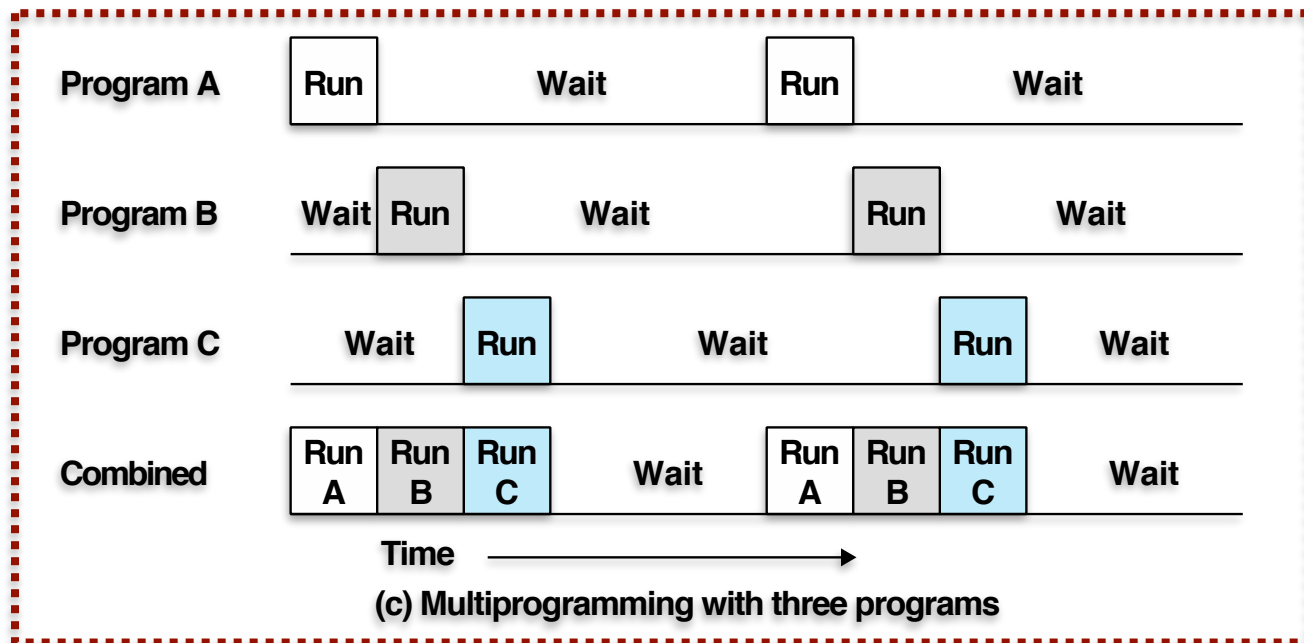
Multiprogramming: Two Programs

- When one job needs to wait for I/O, the processor can switch to the other job, which is likely not waiting for I/O



Multiprogramming: Three Programs


- Memory is expanded to hold three, four, or more programs and switch among all of them
- A.k.a. **multitasking**



Multiprogramming: Example

	JOB1	JOB2	JOB3
Type of job	Heavy compute	Heavy I/O	Heavy I/O
Duration	5 min	15 min	10 min
Memory required	50 M	100 M	75 M
Need disk?	No	No	Yes
Need terminal?	No	Yes	No
Need printer?	No	No	Yes

Multiprogramming: Effects on Resource Utilisation



Time-Sharing Systems

- Can be used to **handle multiple interactive jobs**
- Processor time is shared among multiple users
- Principle objective is to **minimise response time**
- Multiple users simultaneously access the system through terminals, with the OS **interleaving** the execution of each user program in **a short burst or quantum of computation**



**Time
Slicing**

Compatible Time-Sharing Systems







CTSS

- One of the first time-sharing operating systems
- Developed at MIT by a group known as Project MAC
- Ran on a computer with 32,000 36-bit words of main memory, with the resident monitor consuming 5000 of that
- To simplify both the monitor and memory management, a program was always loaded to start at the location of the 5000th word

Time Slicing

- System clock generates **interrupts** at a rate of approximately one every 0.2 seconds
- At each interrupt, OS regained control and could assign processor to another user
- At regular **time intervals**, the current user would be preempted and another user loaded in
- Old user programs and data were written out to disk; before new user programs and data were read in

CTSS: Operations



What are the major advances in OS development?

Major Achievements

- Operating systems are among the most complex pieces of software ever developed

Major advances in OS development:

- Processes
- Memory management
- Information protection and security
- Scheduling and resource management
- System structure

The Concept of Process

- Fundamental to the structure of operating system

A *process* can be defined as:

a program in execution

an instance of a running program

the entity that can be assigned to, and executed on, a processor

a unit of activity characterized by a single sequential thread of execution, a current state, and an associated set of system resources

Development of the Concept of Process

- Three major lines of computer system development that created problems in **timing and synchronisation**
- **Multiprogramming batch operation:**
 - Processor is switched among the various programs residing in main memory
- **Time sharing:**
 - Be responsive to the individual users but be able to support many users simultaneously
- **Real-time transaction systems:**
 - A number of users are entering queries or updates against a database

Process: Three Components

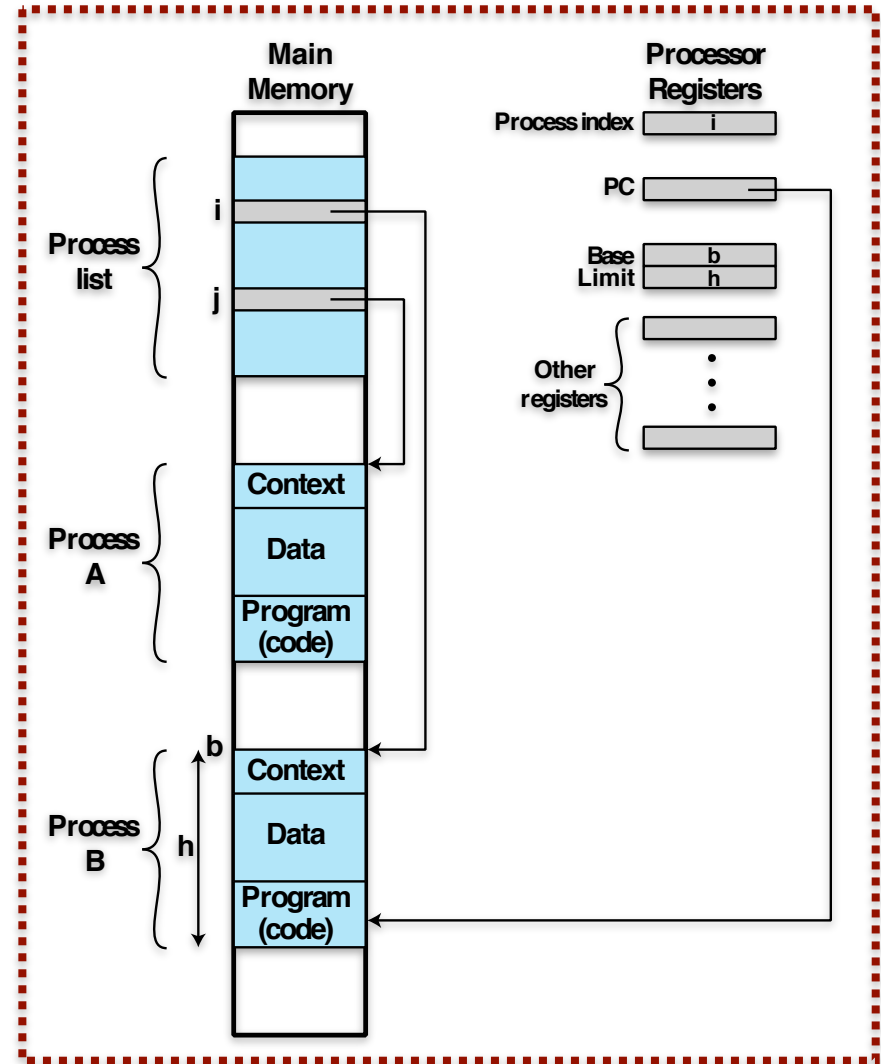
- A process contains three components:

- An executable program
- Associated data needed by the program (e.g. variables, work space, buffers, etc.)
- Execution context (or process state) of the program

- The **execution context** is essential:
 - It is the internal data by which the OS is able to **supervise** and **control** the process
 - Includes the contents of the various process registers
 - Includes information such as the priority of the process; and whether the process is waiting for the completion of a particular I/O event

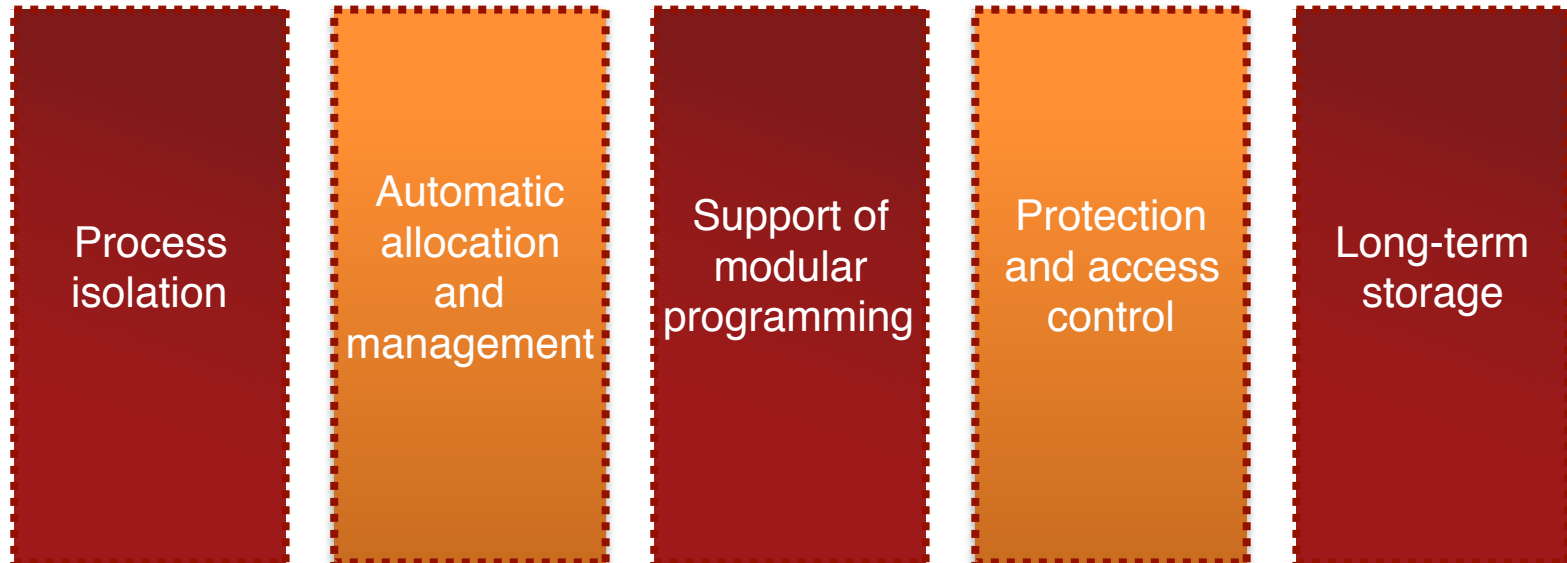
Process Management

- The process is realised as a data structure.
- The entire **state of the process** at any instant is contained in **its context**.
- New features can be designed and incorporated into the OS — by expanding the context to include any new information needed to support the features.



Memory Management

- OS has five principal storage management responsibilities:



- Can be achieved using **virtual memory** and **file system facilities**

Information Protection and Security

- The nature of the **threat** that concerns an organisation will vary greatly depending on the circumstances
- The concern involves **controlling access** to computer systems and the information stored in the systems
- **Main issues:**
 - Availability
 - Confidentiality
 - Data integrity
 - Authenticity

Scheduling and Resource Management

- The key responsibility of an OS is managing various system resources (main memory, I/O devices, processors)
- Resource allocation policies must consider:
 - **Fairness**: equal and fair access to resources by competing processes
 - **Differential responsiveness**: discriminate among different classes of processes
 - **Efficiency**: maximise throughput and minimise response time; accomodate as many users as possible

OS Architecture

- Demands on operating systems require new approaches to organising the OS architecture

Different approaches and design elements have been tried:

- microkernel architecture
- multithreading
- symmetric multiprocessing
- distributed operating systems
- object-oriented design



MONASH
University

The OS architecture for Windows, Unix/Linux, and Android


Traditional Unix Systems

- Developed at Bell Labs and became operational on a PDP-7 in 1970
 - Incorporated many ideas from **Multics** (and also **CTSS**)
 - PDP-11 was a milestone because it first showed that UNIX would be an OS for all computers
- Next milestone was rewriting **UNIX in the programming language C**
 - Demonstrated the advantages of using a high-level language for system code
- Described in a technical journal for the first time in 1974
- First widely available version outside Bell Labs was Version 6 in 1976
 - *Version 7* released in 1978 is the ancestor of most modern UNIX systems
- Commercially marketed as Unix System III and **Unix System V**
- Most important of the non-AT&T systems was **Unix BSD**

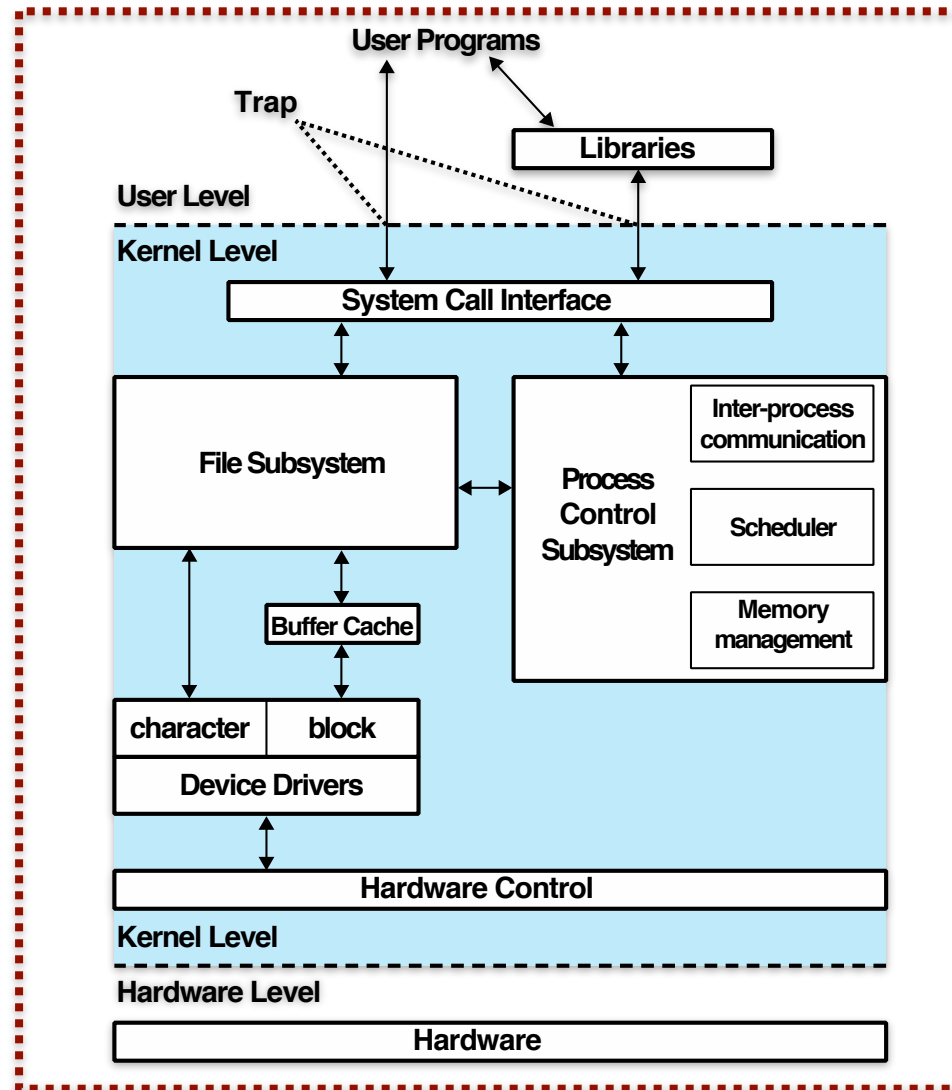


Berkeley
Software
Distribution

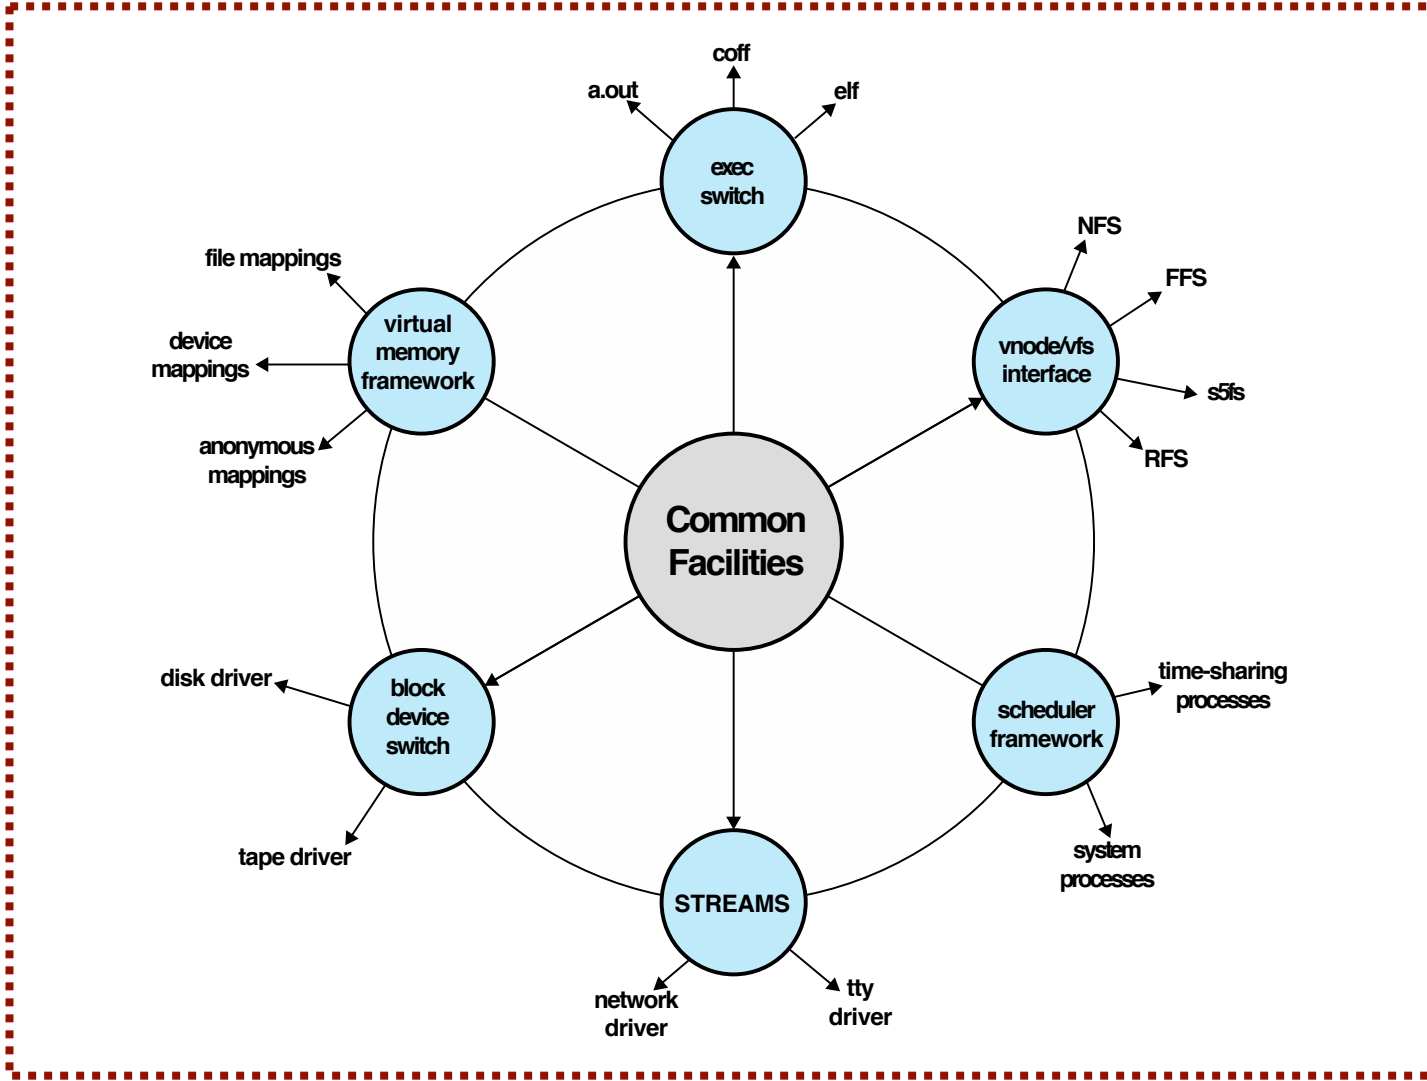
Unix: General Architecture



Unix Kernel: Traditional Architecture



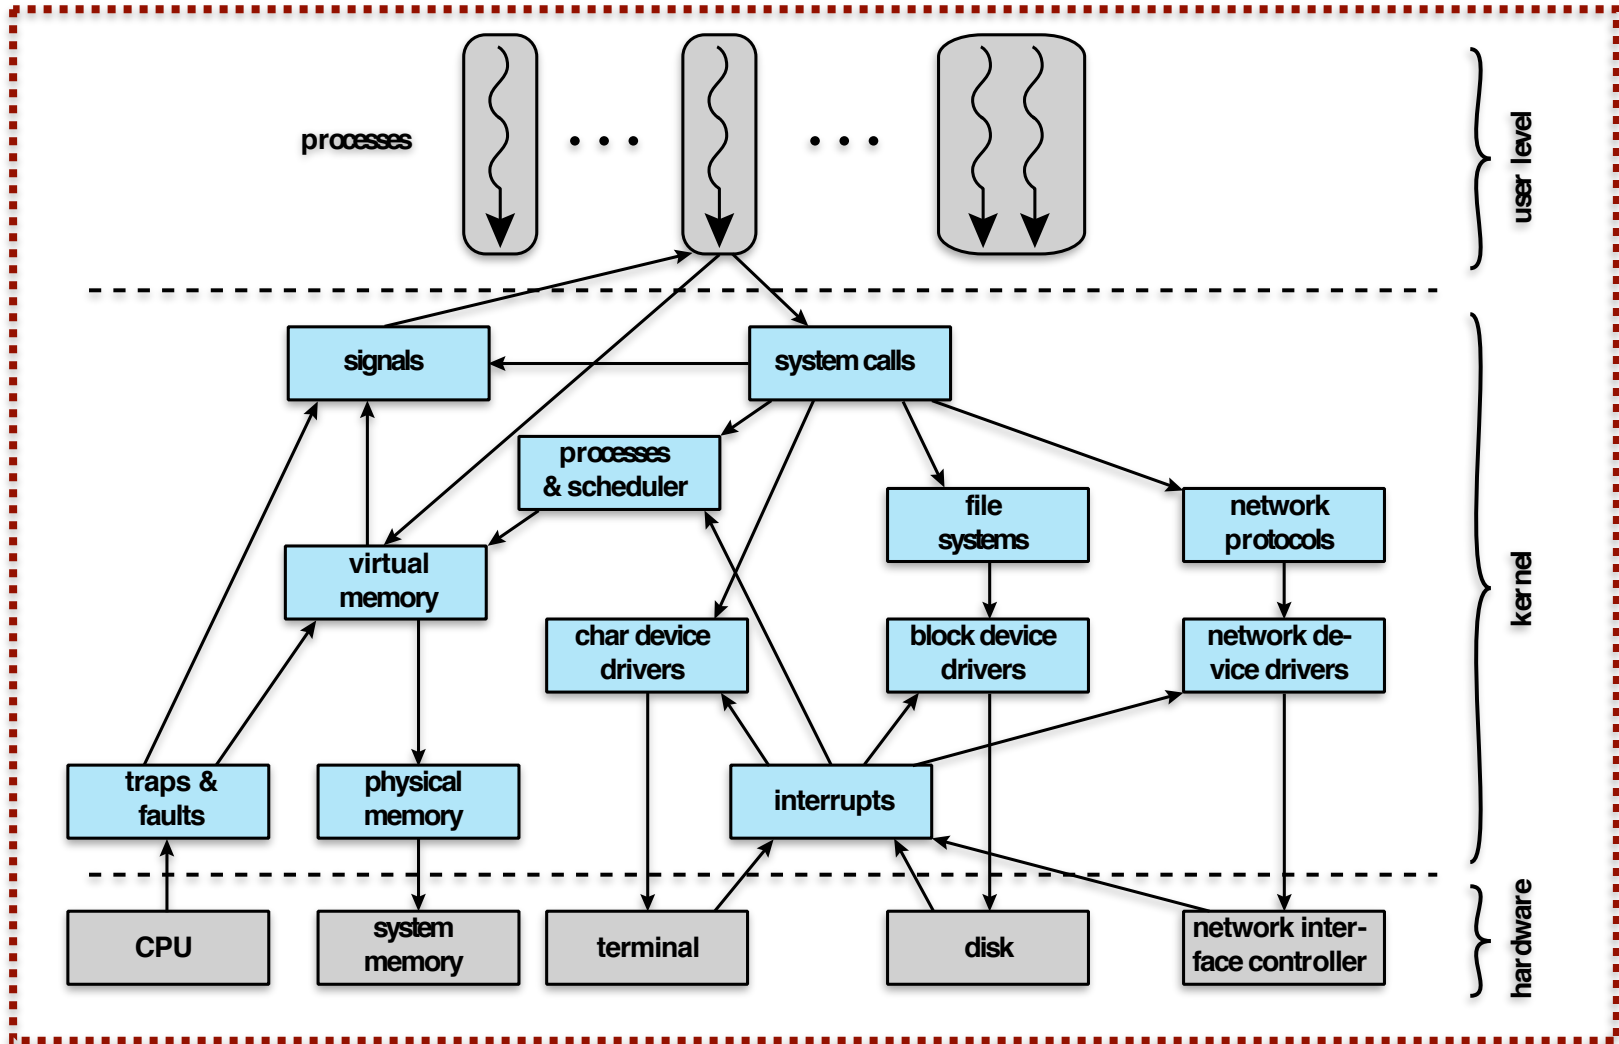
Unix Kernel: Modern Architecture



Linux Systems

- Started out as a UNIX variant for the IBM PC
- Initial version was written by **Linus Torvalds**, a Finnish student of computer science
- Linux was first posted on the Internet in 1991
- Today it is a full-featured UNIX system that runs on several platforms
- Free and source code is available (**GNU Public License**)
- Key to success has been the availability of free software packages
- **Highly modular** and easily configured

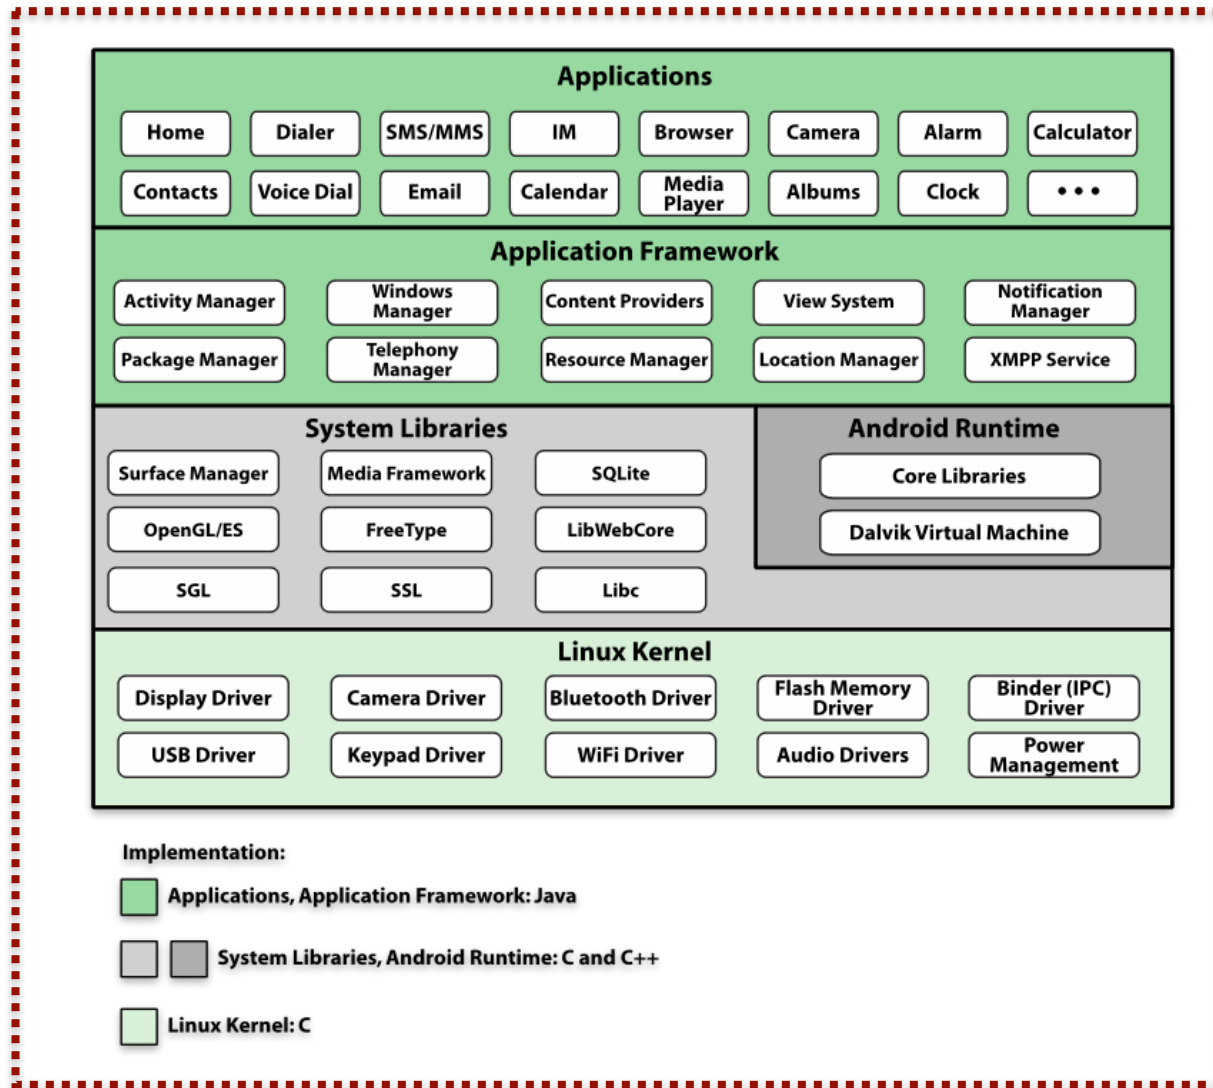
Linux: Kernel Components



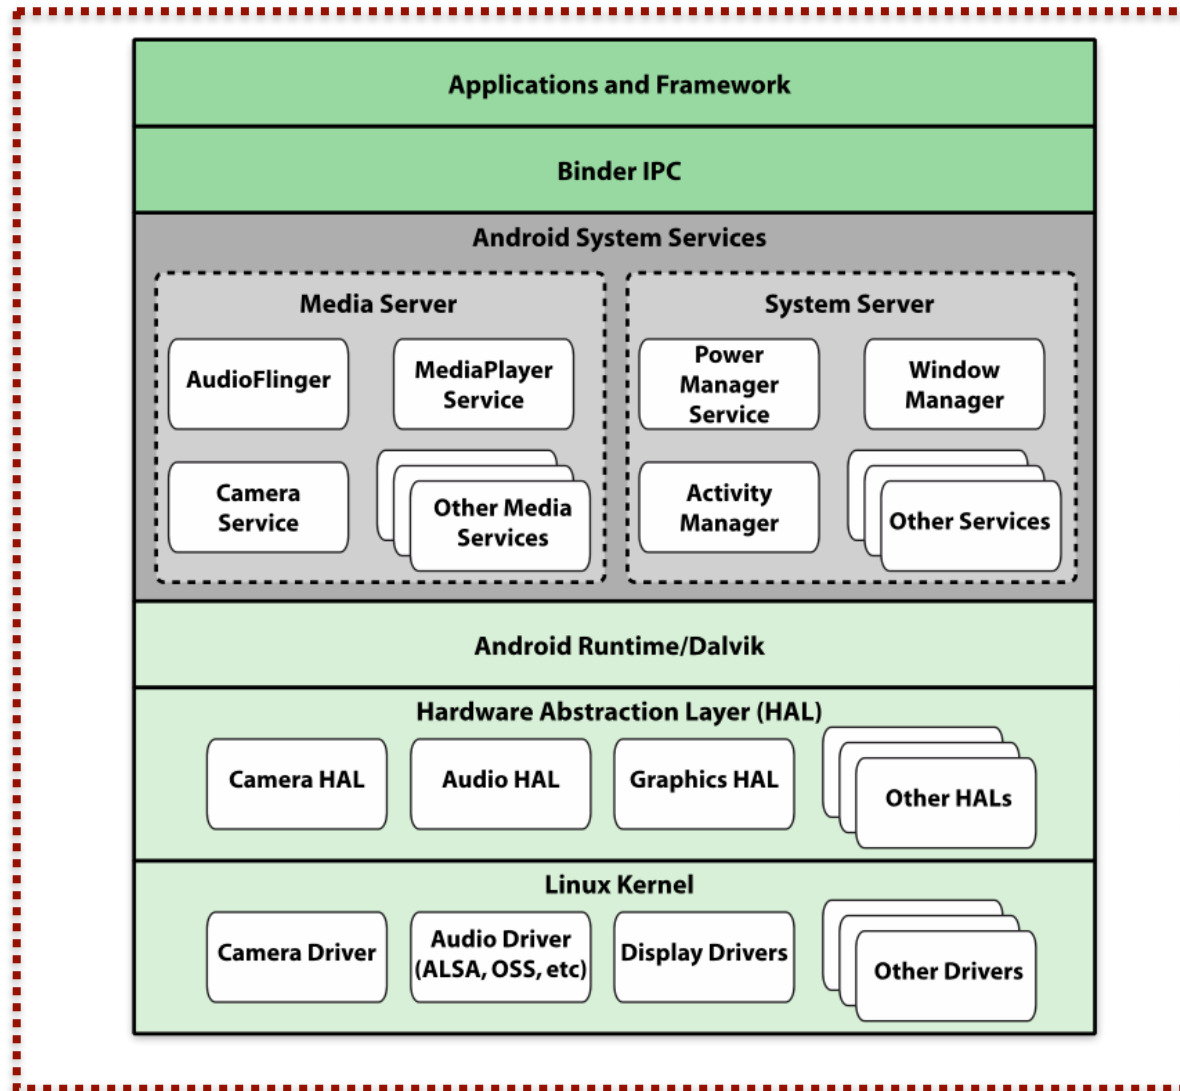
Android OS

- A Linux-based system originally designed for touchscreen mobile devices such as smartphones and tablet computers
- Development was done by **Android Inc.**, which was bought by Google in 2005
- 1st commercial version (Android 1.0) was released in 2008
- Most recent version is Android 7.0 (Nougat)
- The **Open Handset Alliance (OHA)** was responsible for the Android OS releases as an open platform
- The open-source nature of Android has been the key to its success

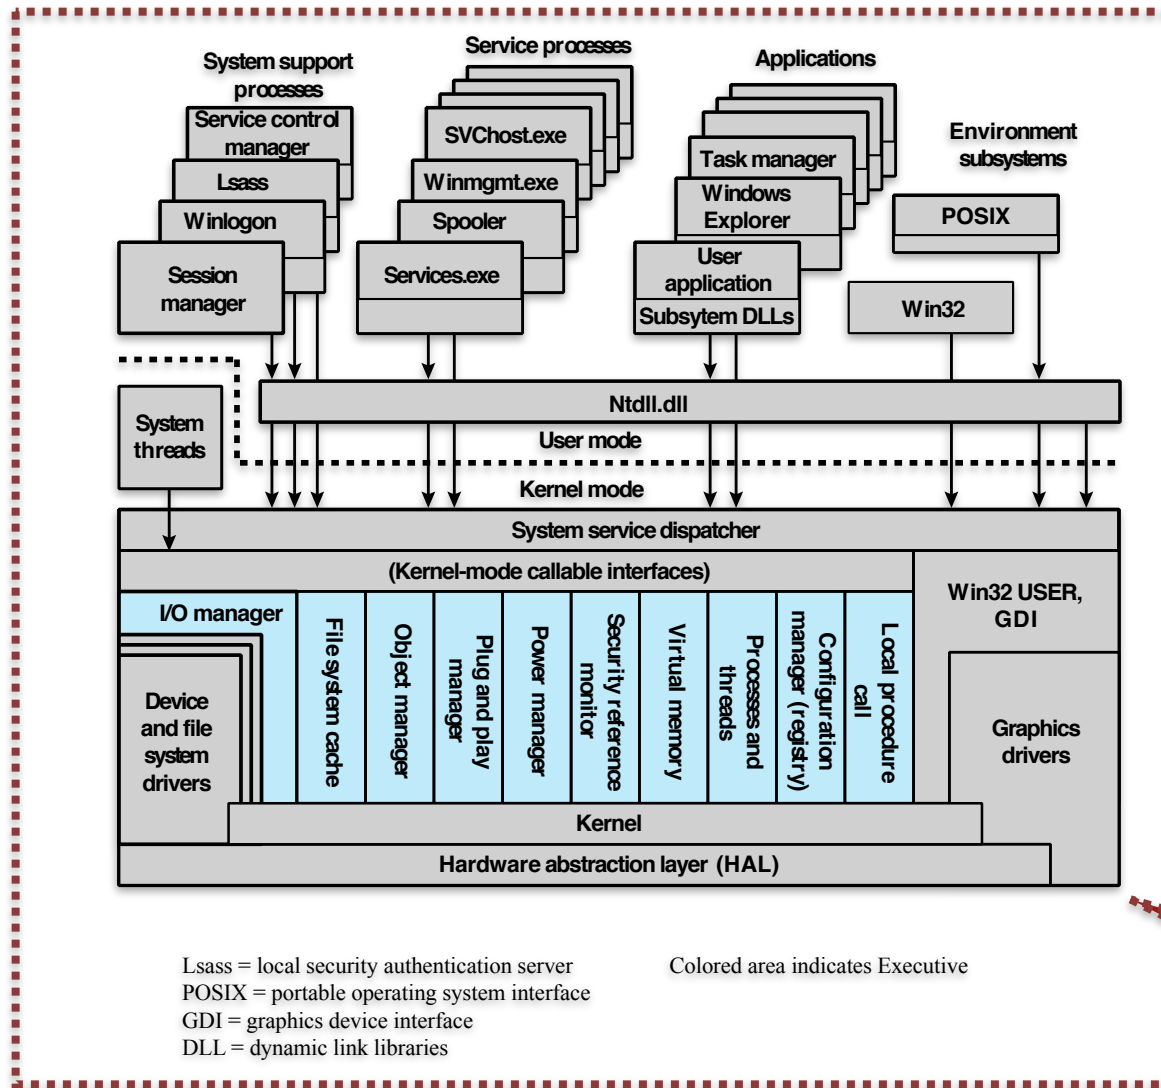
Android: Software Architecture



Android: System Architecture



Microsoft Windows: System Architecture



Windows 8

Summary

- So far, we have discussed:
 - Objectives and functions of OS
 - Evolution of operating systems
 - Key design areas of modern OS
 - OS architecture of Unix/Linux, Android, and Windows
- Next week:
 - Concepts of Processes and Threads

Reminder: The first practical this week.