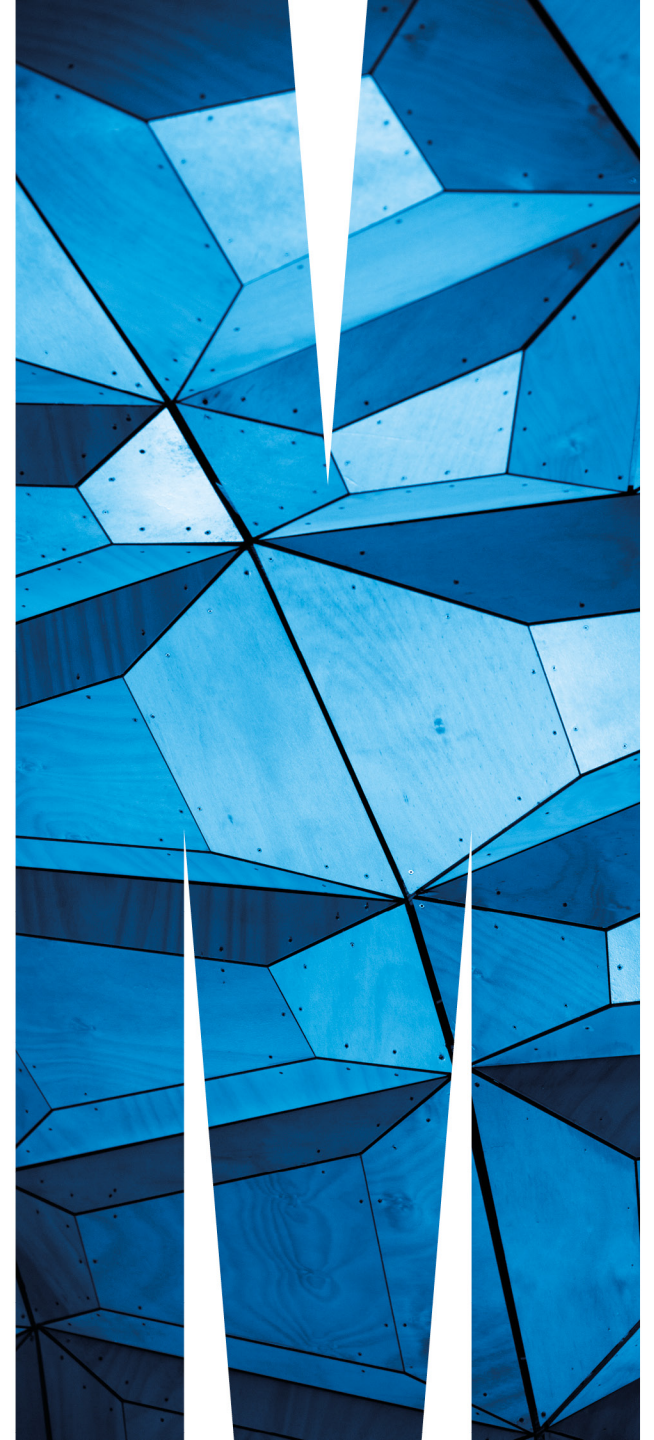


# FIT2100 Semester 2 2017

## Lecture 4 (Part B): Concurrency (Part 1)

### (Reading: Stallings, Chapter 5)

Jojo Wong



# Lecture 4 (Part 2): Learning Outcomes

- Upon the completion of this lecture, you should be able to:
  - Discuss the basic concepts related to **concurrency**
  - Understand the concept of **race condition**
  - Describe the **mutual exclusion** requirements

Why concurrency is important in supporting multiprocessing?

# Multiple Processes

- OS design is concerned with the management of processes and threads:
  - Multiprogramming
  - Multiprocessing
  - Distributed processing

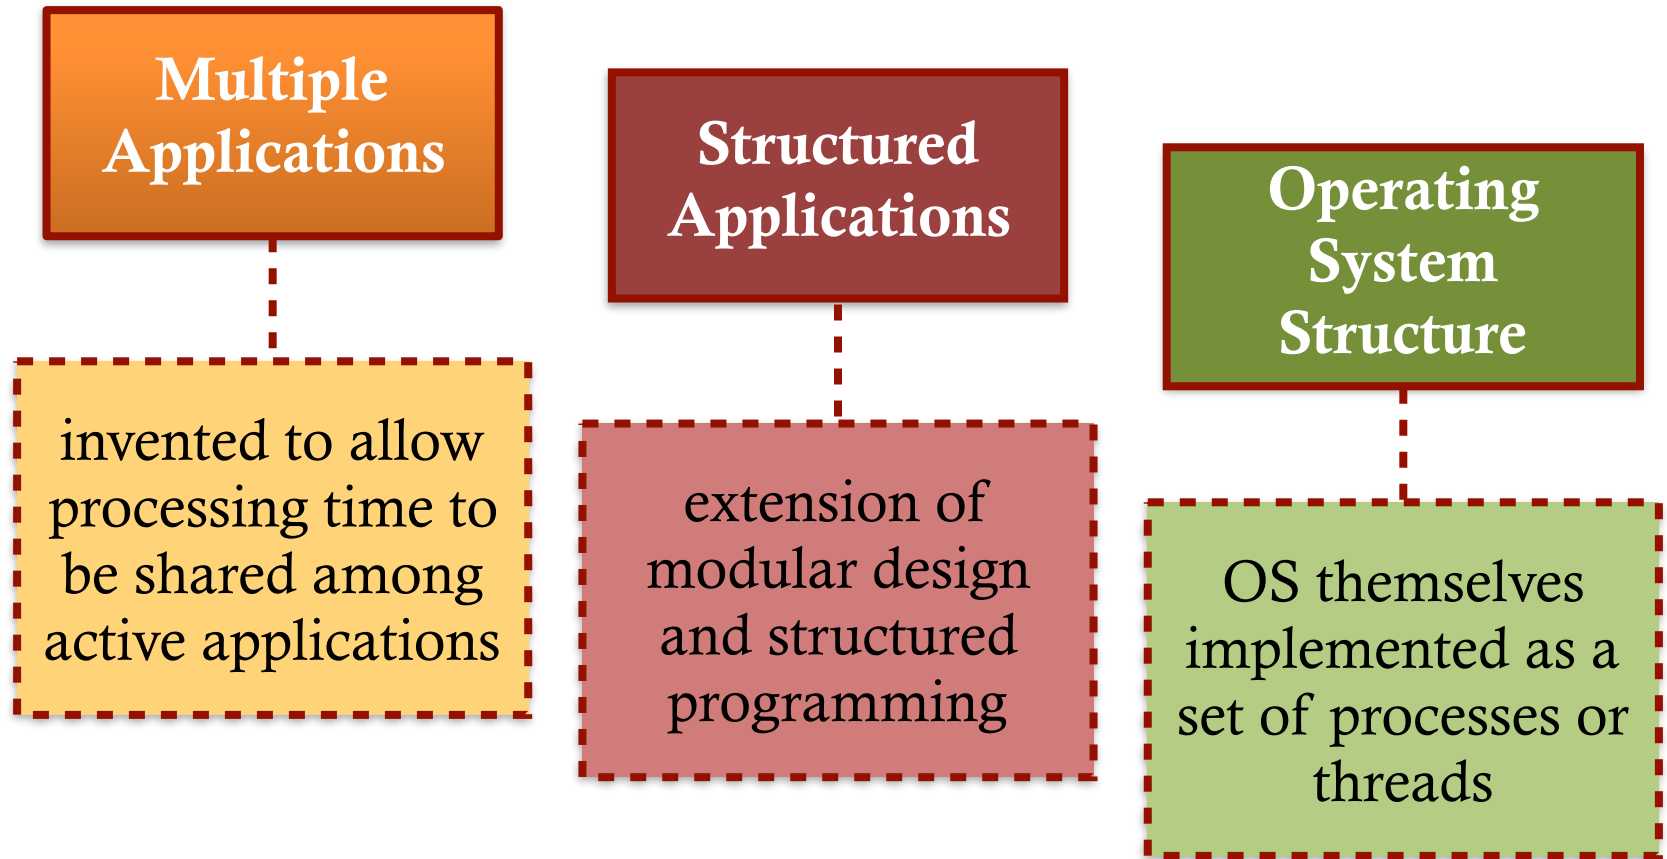
mgt of multiple processes within a uniprocessor system

The diagram consists of three red boxes with dashed borders. The top box is connected to 'Multiprogramming' by a line. The middle box is connected to 'Multiprocessing' by a line. The bottom box is connected to 'Distributed processing' by a line. The boxes are arranged in a descending staircase pattern from top-right to bottom-left.

mgt of multiple processes within a multi-processor system

mgt of multiple processes on a multiple distributed computer system

# Concurrency: Three Different Contexts





# Concurrency: Terminology

<b>atomic operation</b>	A function or action implemented as a sequence of one or more instructions that appears to be indivisible; that is, no other process can see an intermediate state or interrupt the operation. The sequence of instruction is guaranteed to execute as a group, or not execute at all, having no visible effect on system state. Atomicity guarantees isolation from concurrent processes.
<b>critical section</b>	A section of code within a process that requires access to shared resources and that must not be executed while another process is in a corresponding section of code.
<b>deadlock</b>	A situation in which two or more processes are unable to proceed because each is waiting for one of the others to do something.
<b>livelock</b>	A situation in which two or more processes continuously change their states in response to changes in the other process(es) without doing any useful work.
<b>mutual exclusion</b>	The requirement that when one process is in a critical section that accesses shared resources, no other process may be in a critical section that accesses any of those shared resources.
<b>race condition</b>	A situation in which multiple threads or processes read and write a shared data item and the final result depends on the relative timing of their execution.
<b>starvation</b>	A situation in which a runnable process is overlooked indefinitely by the scheduler; although it is able to proceed, it is never chosen.

# Concurrency: Principles

- **Uniprocessor:** the relative speed of execution of processes cannot be predicted
  - depends on activities of other processes
  - the way OS handles interrupts
  - scheduling policies of the OS
- **Multiprocessor:** *interleaving and overlapping*
  - can be viewed as examples of concurrent processing
  - both present the same problems

# Concurrency: Difficulties

- Sharing of global resources  **race condition**
- Difficult for OS to manage the allocation of resources optimally  **deadlock**
- Difficult to locate programming errors — results are not deterministic and reproducible



What are the concurrency problems?

# Concurrency: Problems

- Concurrent access to shared data may result in data inconsistency — **Race Condition**
- A problem exists in multiprogramming on uni- and multi-processors
- Maintaining **data consistency** requires mechanisms to ensure the orderly execution of cooperating process

# Race Condition

- Occurs when multiple processes or threads read and write data items
- Final result depends on the **order of execution**
  - “**Loser**” of the race is the process that updates last and will determine the final value of the variable
- To prevent race conditions, concurrent processes must be **synchronised**

# Concurrency Problem: Uniprocessor Multiprogramming

- Consider the following procedure:

```
void echo()  
{  
    char-in = getchar();  
    char-out = char-in;  
    putchar(char-out);  
}
```



**Multiprogramming  
with uniprocessor**

- Read a character from the keyboard and store in **char-in**.
- Transfer to **char-out** before being sent for display.
- Consider two different applications — **P1** and **P2** — make a call to this procedure.

# Concurrency Problem: Uniprocessor Multiprogramming

- Consider the following procedure:

```
void echo()  
{  
    char-in = getchar();  
    char-out = char-in;  
    putchar(char-out);  
}
```



Multiprogramming  
with uniprocessor

- P1 invokes `echo()` and is interrupted immediately after `getchar` returns its value and stores it in `char-in` (e.g. `x`).
- P2 is activated and invokes `echo()` and read a char (e.g. `y`) and runs to completion of the procedure.
- When P1 resumes the value of `x` has been overwritten in `char-in` by process P2 and therefore its value `x` is lost.

# Concurrency Problem: Uniprocessor Multiprogramming

- Process P1:

```
char-in = getchar();
```

```
char-out = char-in;  
putchar(char-out);
```

- Process P2:

```
char-in = getchar();  
char-out = char-in;  
putchar(char-out);
```



**TIME**

# Concurrency Problem: Uniprocessor Multiprogramming

- Consider the following procedure:

```
void echo()  
{  
    char-in = getchar();  
    char-out = char-in;  
    putchar(char-out);  
}
```



Multiprogramming  
with uniprocessor

- Assume **only one process at a time** to invoke and be in the `echo` procedure.
- P1 invokes `echo()` and is interrupted immediately after `getchar` returns its value and stores it in `char-in (x)`.
- P2 is activated and invokes `echo()`.
- But since P1 is still inside the procedure, and currently suspended — **P2 is blocked** from entering the procedure.

# Concurrency Problem (Summary):

- P1 invokes the **echo** procedure and is interrupted immediately after **getchar** returns its value and stores it in **char-in** (e.g. **x**).
- P2 is activated and invokes **echo** procedure and since the **echo** procedure is used by process P1, **P2 is blocked from further execution**.
- At some later time, P1 is resumed and completes the execution of **echo** and the proper input character will be displayed.
- When P1 exits **echo**, **this removes the block on P2**.
- When P2 is later resumed, the **echo** procedure is successfully invoked.





MONASH  
University

How about concurrency problems  
with multiprocessors?

# Concurrency Problems: Multiprocessor Multiprogramming

- Same problem arises even when the processes — P1 and P2 — runs on **different processors** accessing unprotected shared variables.
- The solution outlined in the previous slides can work here.
- **Protecting and controlling access to shared resources are critical.**

## Question: Race Condition

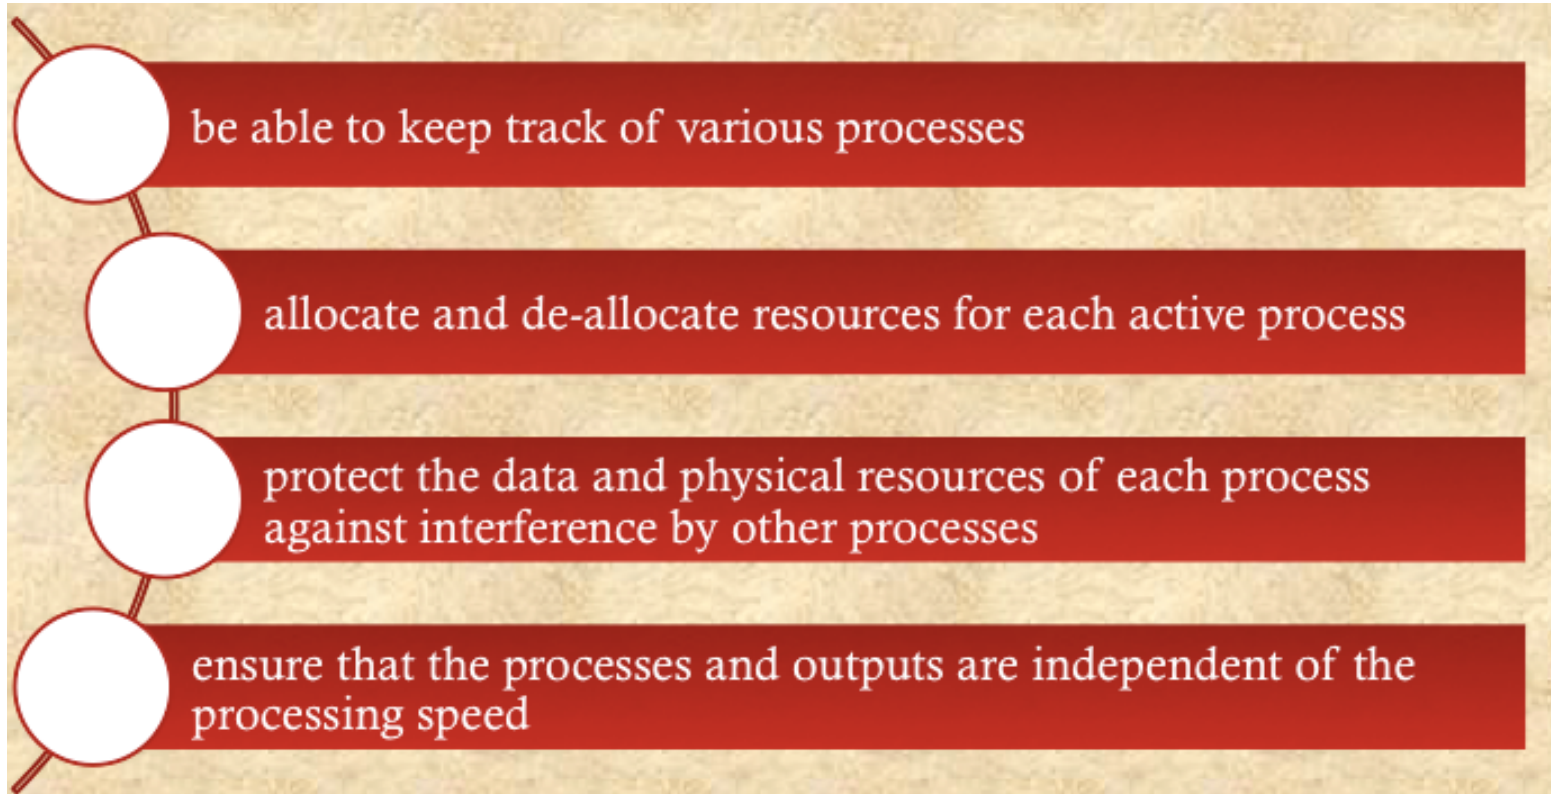
- Assume P1 and P2 share two variables **a** and **b** with initial values of **a** = 1 and **b** = 2
- P1 executes the statement: **a** = **a** + **b**
- P2 executes the statement: **b** = **a** + **b**
- What values are **a** and **b** if P1 executes before P2?
- What values are **a** and **b** if P2 executes before P1?



MONASH  
University

# What are the responsibilities of OS?

# Operating System: Concerns



How do processes interact with each other?

# Types of Process Interaction

Degree of Awareness	Relationship	Influence that One Process Has on the Other	Potential Control Problems
Processes unaware of each other	Competition	<ul style="list-style-type: none"> <li>•Results of one process independent of the action of others</li> <li>•Timing of process may be affected</li> </ul>	<ul style="list-style-type: none"> <li>•Mutual exclusion</li> <li>•Deadlock (renewable resource)</li> <li>•Starvation</li> </ul>
Processes indirectly aware of each other (e.g., shared object)	Cooperation by sharing	<ul style="list-style-type: none"> <li>•Results of one process may depend on information obtained from others</li> <li>•Timing of process may be affected</li> </ul>	<ul style="list-style-type: none"> <li>•Mutual exclusion</li> <li>•Deadlock (renewable resource)</li> <li>•Starvation</li> <li>•Data coherence</li> </ul>
Processes directly aware of each other (have communication primitives available to them)	Cooperation by communication	<ul style="list-style-type: none"> <li>•Results of one process may depend on information obtained from others</li> <li>•Timing of process may be affected</li> </ul>	<ul style="list-style-type: none"> <li>•Deadlock (consumable resource)</li> <li>•Starvation</li> </ul>

What are the control problems with concurrent processes?



# Control Problems

- Concurrent processes come into **conflict** when they are competing for the same system resource — I/O devices, memory, processor time, etc.

**In the case of competing processes three control problems must be faced:**

- **mutual exclusion**
- **deadlock**
- **starvation**

# Mutual Exclusion

- Suppose  $n$  processes all **competing** to use some shared data.
- Each process has a code segment — **critical section** — where the shared data is accessed or manipulated.
- Ensure that when one process is executing in its critical section, **no other process is allowed in its critical section.**

# Mutual Exclusion: Example

To enforce mutual exclusion, two functions are provided — **entercritical/exitcritical** with the resource (Ra) as the argument

```
/* PROCESS 1 */  
  
void P1  
{  
    while (true) {  
        /* preceding code */;  
        entercritical (Ra);  
        /* critical section */;  
        exitcritical (Ra);  
        /* following code */;  
    }  
}
```

```
/* PROCESS 2 */  
  
void P2  
{  
    while (true) {  
        /* preceding code */;  
        entercritical (Ra);  
        /* critical section */;  
        exitcritical (Ra);  
        /* following code */;  
    }  
}
```

...

```
/* PROCESS n */  
  
void Pn  
{  
    while (true) {  
        /* preceding code */;  
        entercritical (Ra);  
        /* critical section */;  
        exitcritical (Ra);  
        /* following code */;  
    }  
}
```

# Multiple Shared Data Resources

- The same problem exists even when processes access more than one shared resource.
- Processes must **cooperate** to ensure the shared data are properly managed.
- Control mechanisms are needed to ensure the **integrity** of the shared data.

# Multiple Shared Data Resources: Example

- **P1**

```
a = a + 1 ;  
b = b + 1 ;
```

- **P2**

```
b = 2 * b ;  
a = 2 * a ;
```

- Assuming that **a = b** at the beginning, and consider the following concurrent execution sequence:

```
a = a + 1 ;    /* {P1} */  
b = 2 * b ;    /* {P2} */  
b = b + 1 ;    /* {P1} */  
a = 2 * a ;    /* {P2} */
```

- At the end of this execution, the condition **a = b** *no longer holds!*

# Mutual Exclusion: Requirements

- Mutual Exclusion must be enforced.
- A process that halts must do so without interfering with other processes — **no deadlock or starvation**.
- A process must not be denied access to a critical section when there is no other process is using the shared resources (being manipulated by the critical section code).
- No assumptions are made about relative process speeds or the number of processors.
- A process remains inside its critical section for a finite time only.

# Additional Control Problems

- **Deadlock**: two or more processes are waiting indefinitely for the other processes to release the system resources.
- **Starvation**: indefinite blocking of a process.

# Summary of Lecture 4 (Part 2)

- **Concurrency** is the fundamental concern in supporting multiprogramming, multiprocessing, and distributed processing.
- **Mutual exclusion** is the condition where there is a set of concurrent processes — only one of which is able to access a given resource or perform a given function at any time.

Next week: Concurrency mechanisms,  
deadlocks and starvations.